



DS-2002: Data Systems

Overview of Data Warehouse Systems

Prof. Jon Tupitza



UNIVERSITY *of* VIRGINIA



Modern Data Platform: Solution Scenarios

Big (Unstructured and/or Poly-Schematic) Data Integration and Advanced Analytics

"We want to integrate all our data into our data warehouse"



Modern Data Warehousing

"We're trying to predict which of our customers will churn"



Advanced Analytics

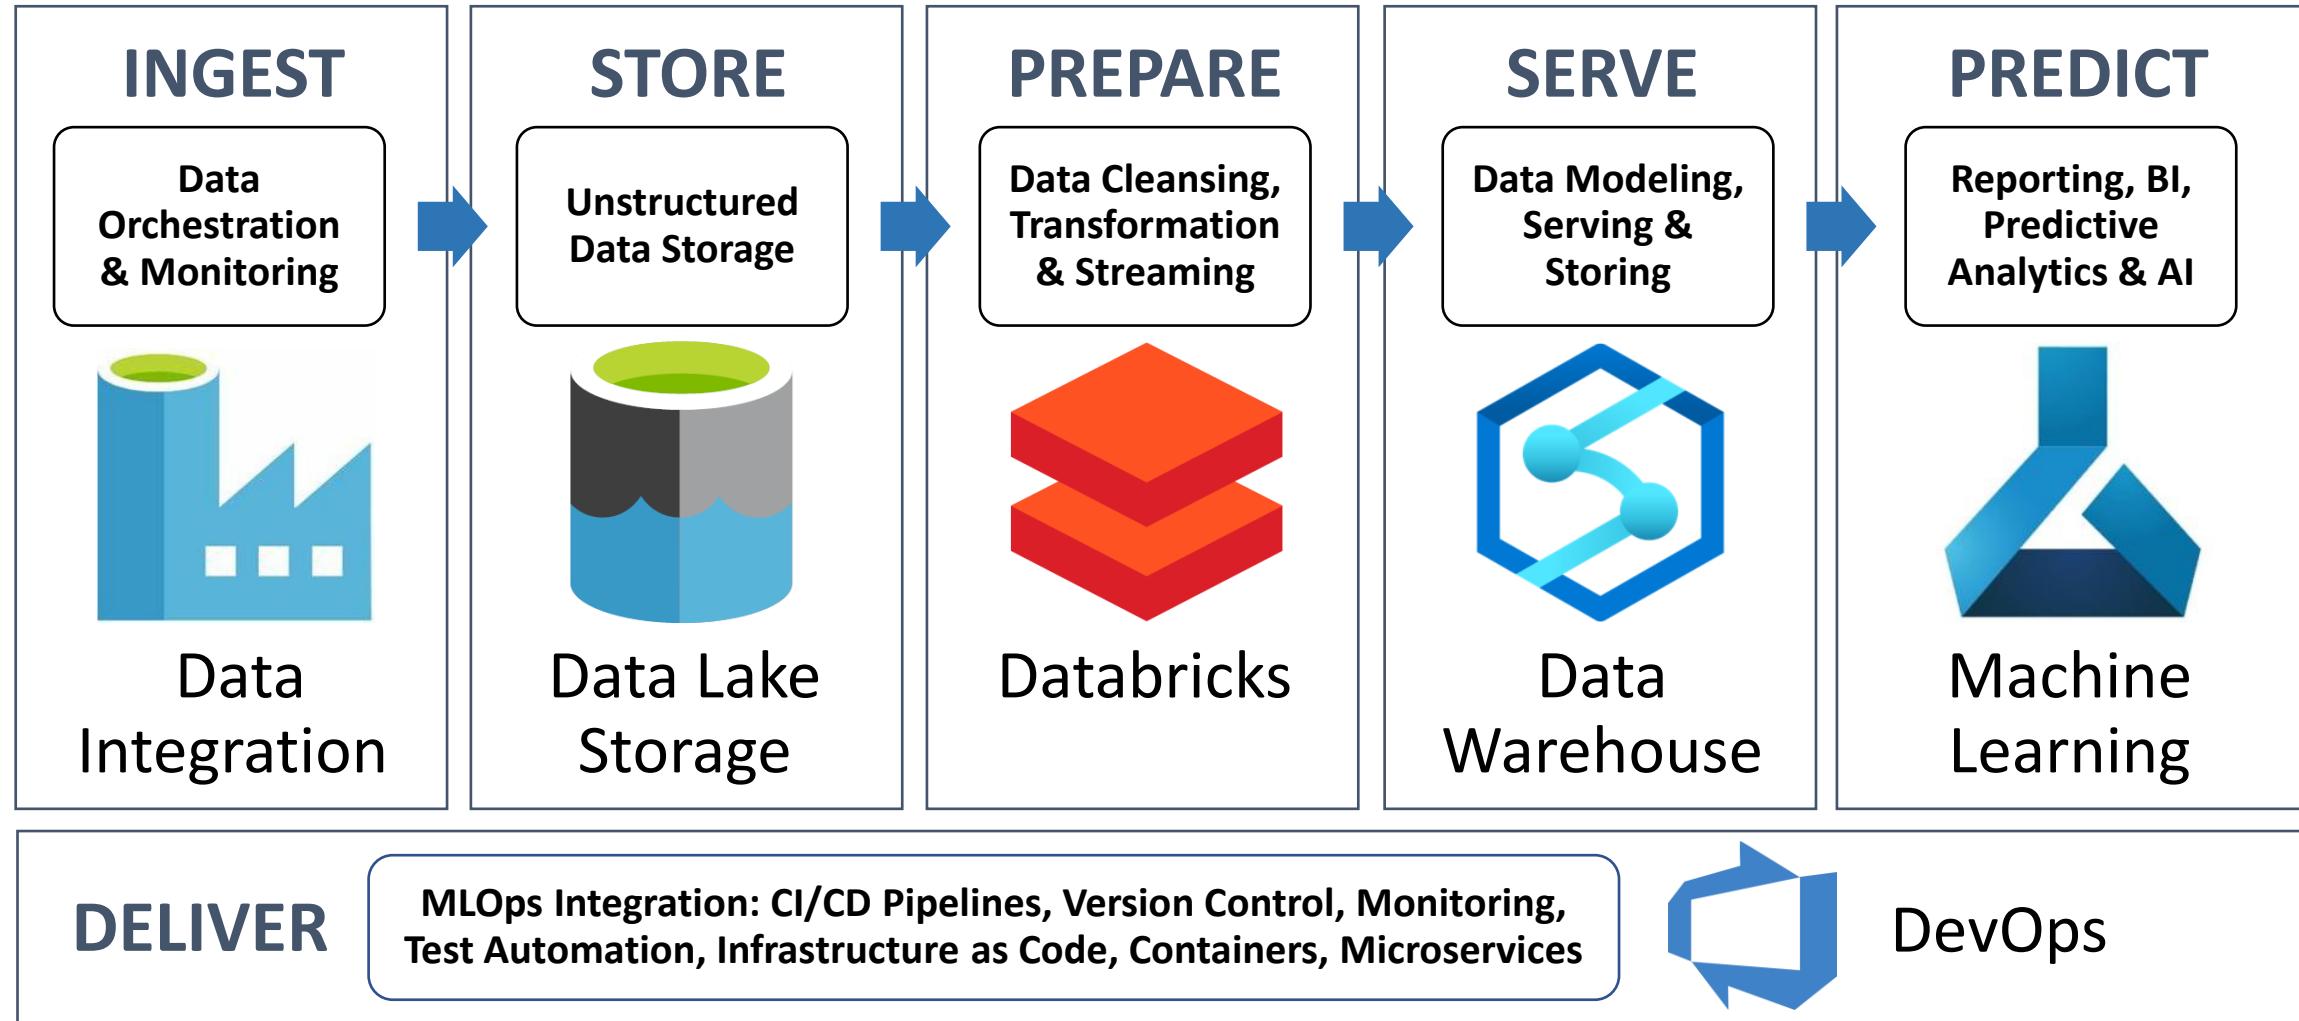
"We're trying to get insights from our devices in real-time"



Real-Time Analytics



Modern Data Platform: Data Services Pipeline



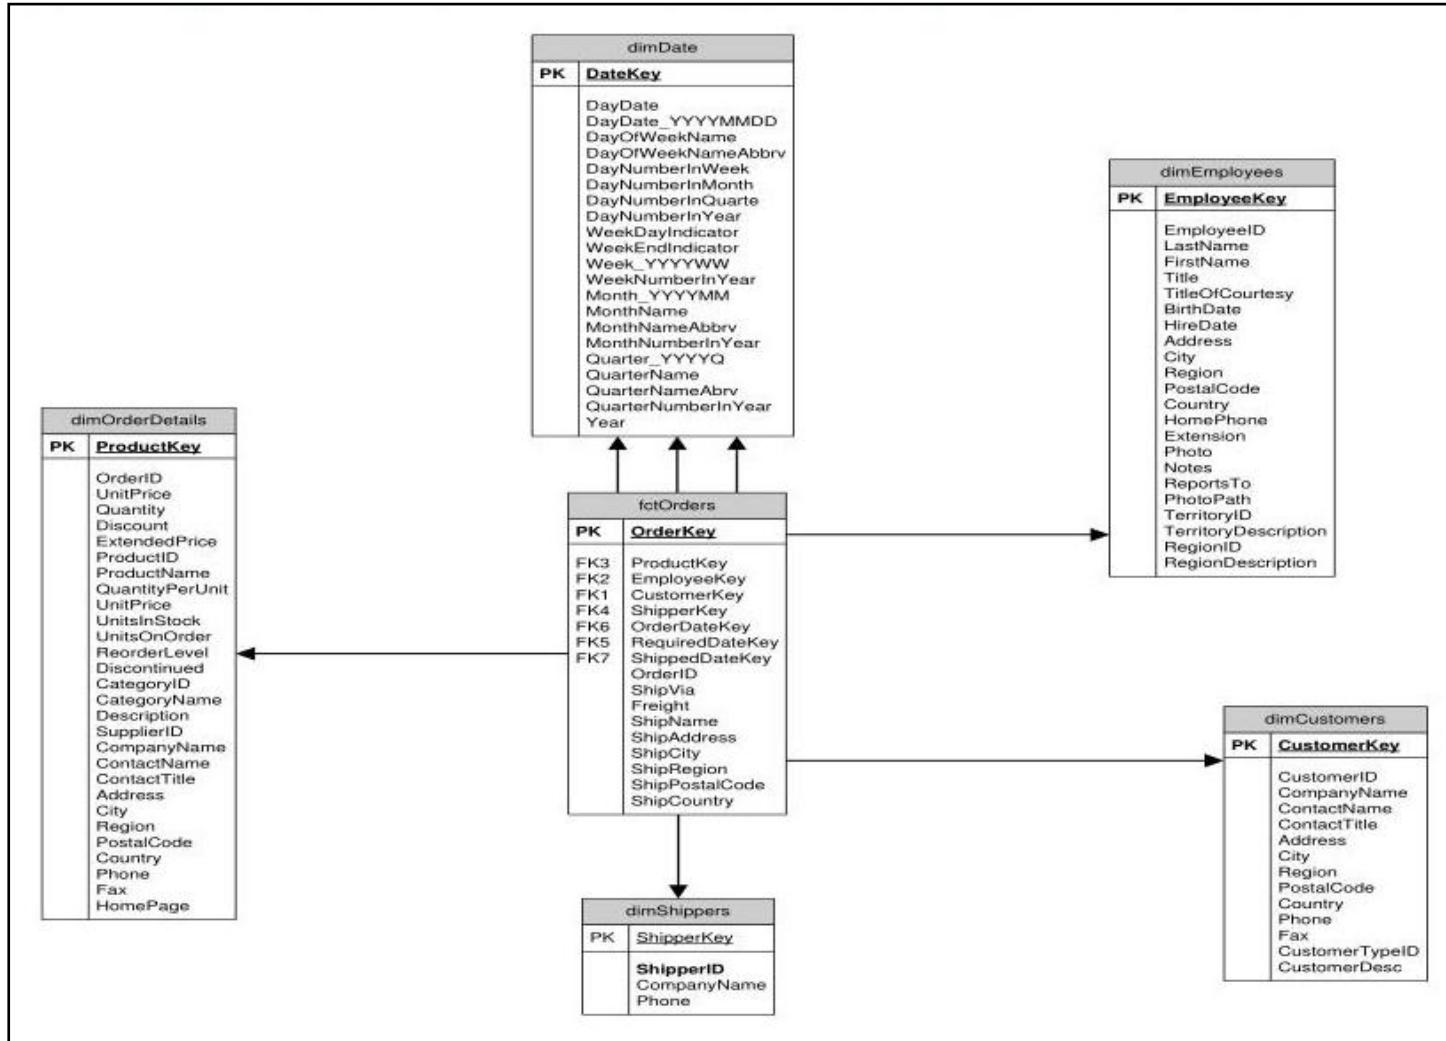
The Data Warehouse Process

How to Approach Designing and Building a Data Warehouse



UNIVERSITY *of* VIRGINIA

OLAP Database: Multi-Dimensional (Star) Schema





The Four-Step Dimensional Design Process

A Time-Honored and Tested Methodology for Delivering Data Marts & Data Warehouses

1. Select the Business Process

- Prioritize Business Processes that Have Significant Business Impact, and that can be Easily Delivered
- Attempting to Boil-the-Ocean Increases the Risk of Project Failure

2. Declare the Grain

- Determine What Level of Detail Must Be Available
- Greater Atomicity Provides Greater Analytic Flexibility
- Aggregated Data Limits Access to Transaction Detail

3. Identify the Dimensions

- Categorical Values Describing Each Business Process
- Defines How Transactions are Viewed or Filtered
- Should Have the Same Meaning Across Processes

4. Identify the Facts

- Numeric Values that Quantify Each Business Process
- Defines How Transactions Are Summarized
- Facts Must Be True-to-the-Grain

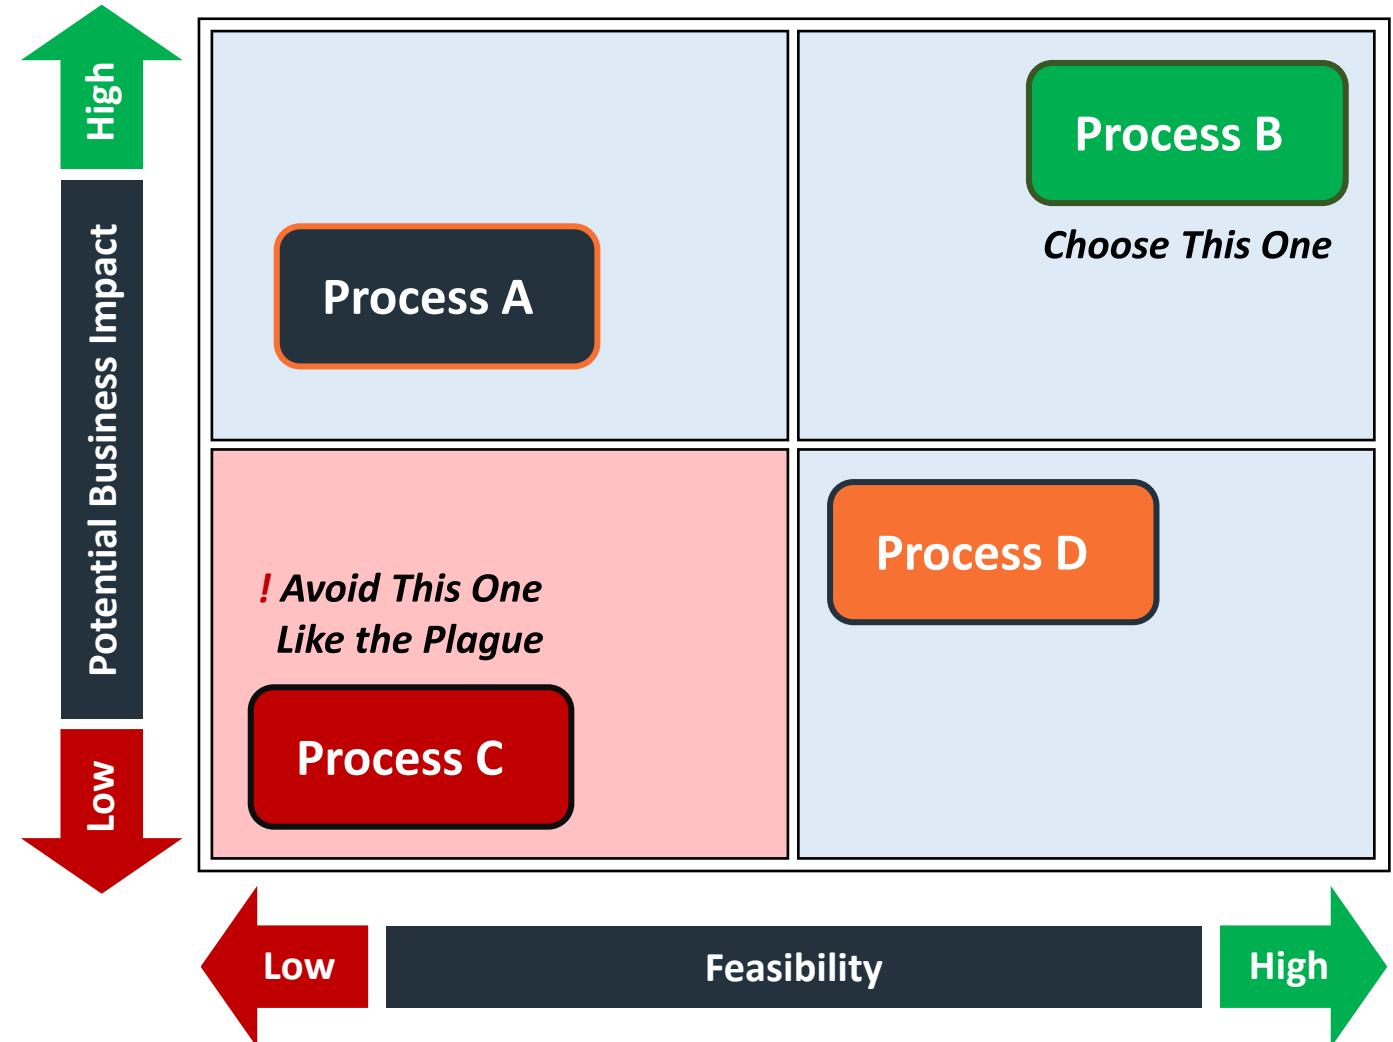
The Data Warehouse Toolkit, by Ralph Kimball



Selecting Business Processes: Prioritizing Requirements

Quadrant Analysis for Prioritizing Requirements:

- Business Process A:**
 - High Potential Business Impact
 - Extremely Difficult to Implement
- Business Process B:**
 - High Potential Business Impact
 - Highly Feasible
- Business Process C:**
 - Very Little Business Impact
 - Extremely Difficult to Implement
- Business Process D:**
 - Little Business Impact
 - Highly Feasible

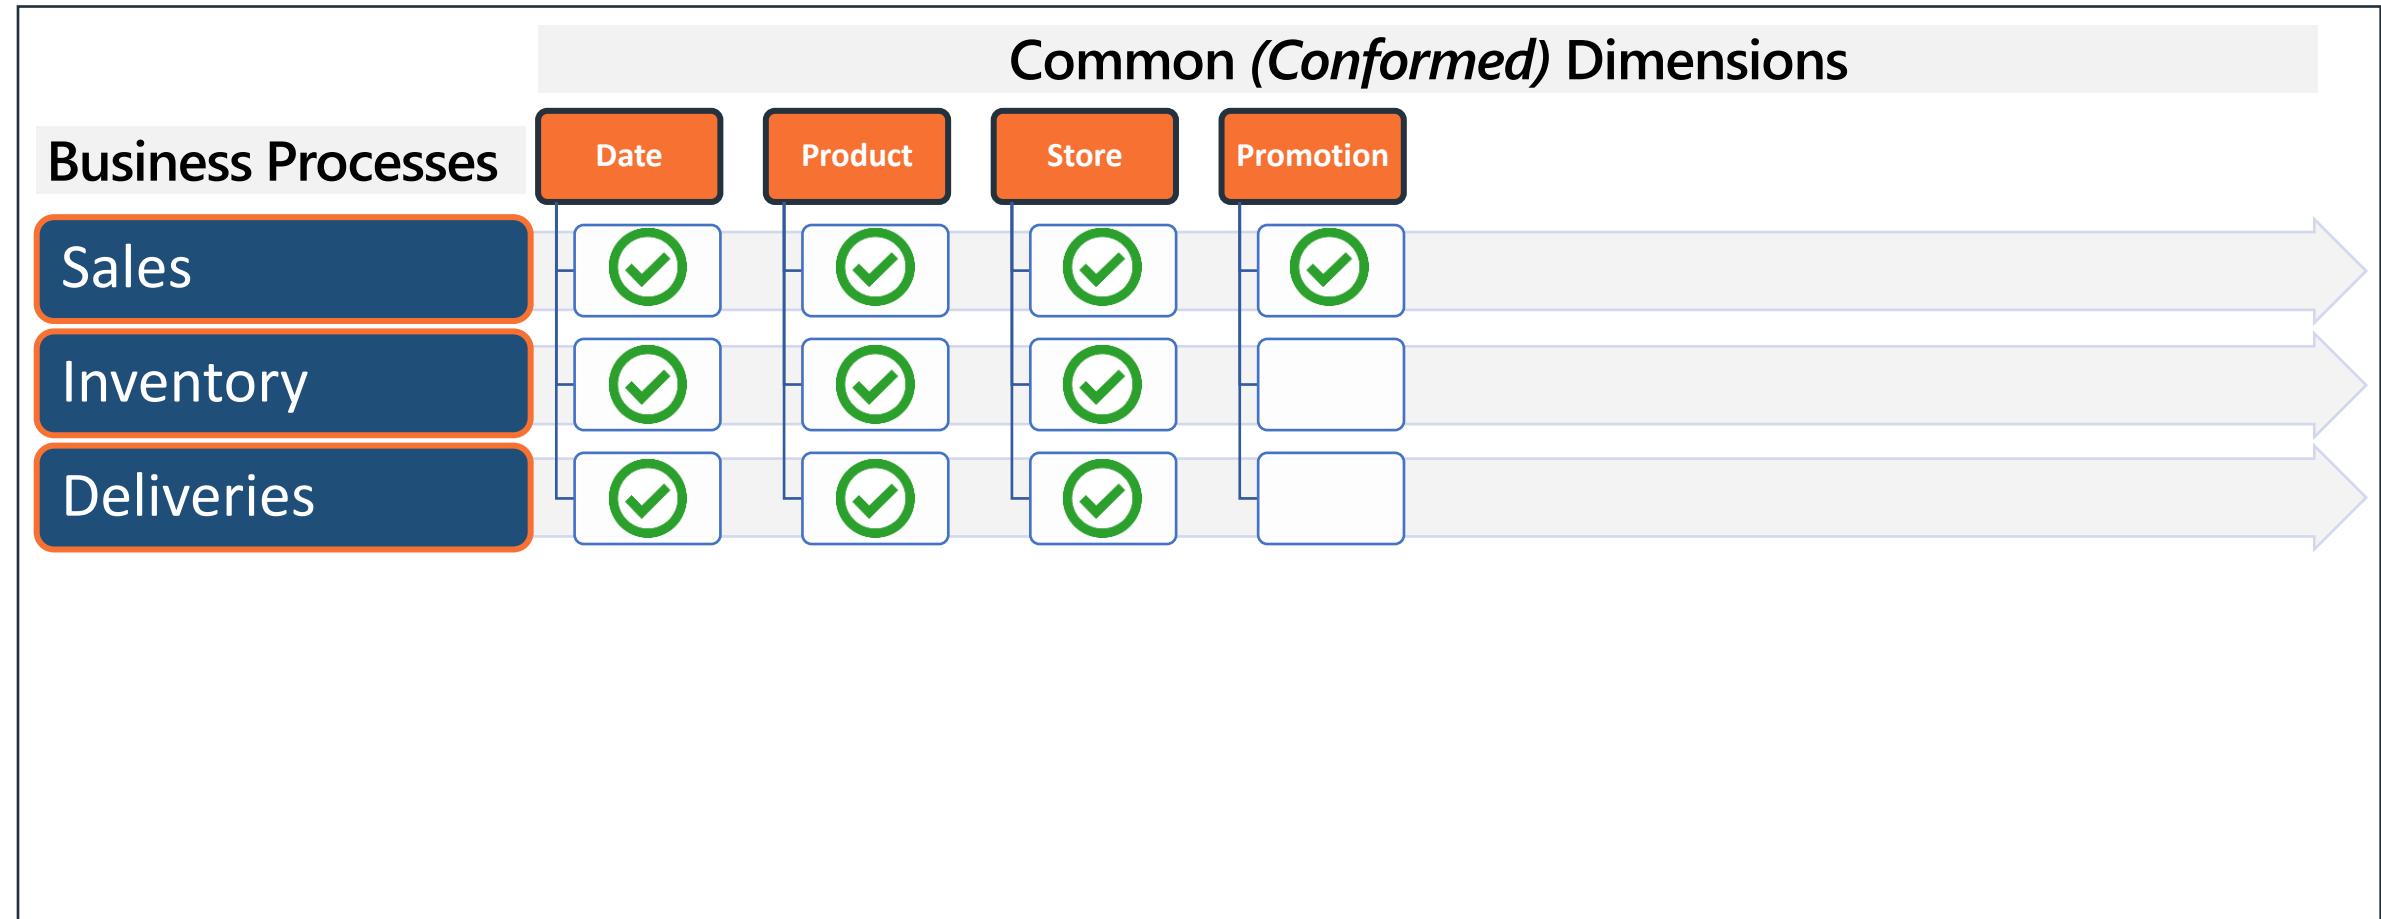


[The Data Warehouse Toolkit](#), by Ralph Kimball



Identifying Dimensions: Data Warehouse Bus Matrix

Using the Same Dimensions Across Multiple Business Processes Enforces a Unified View of the Truth

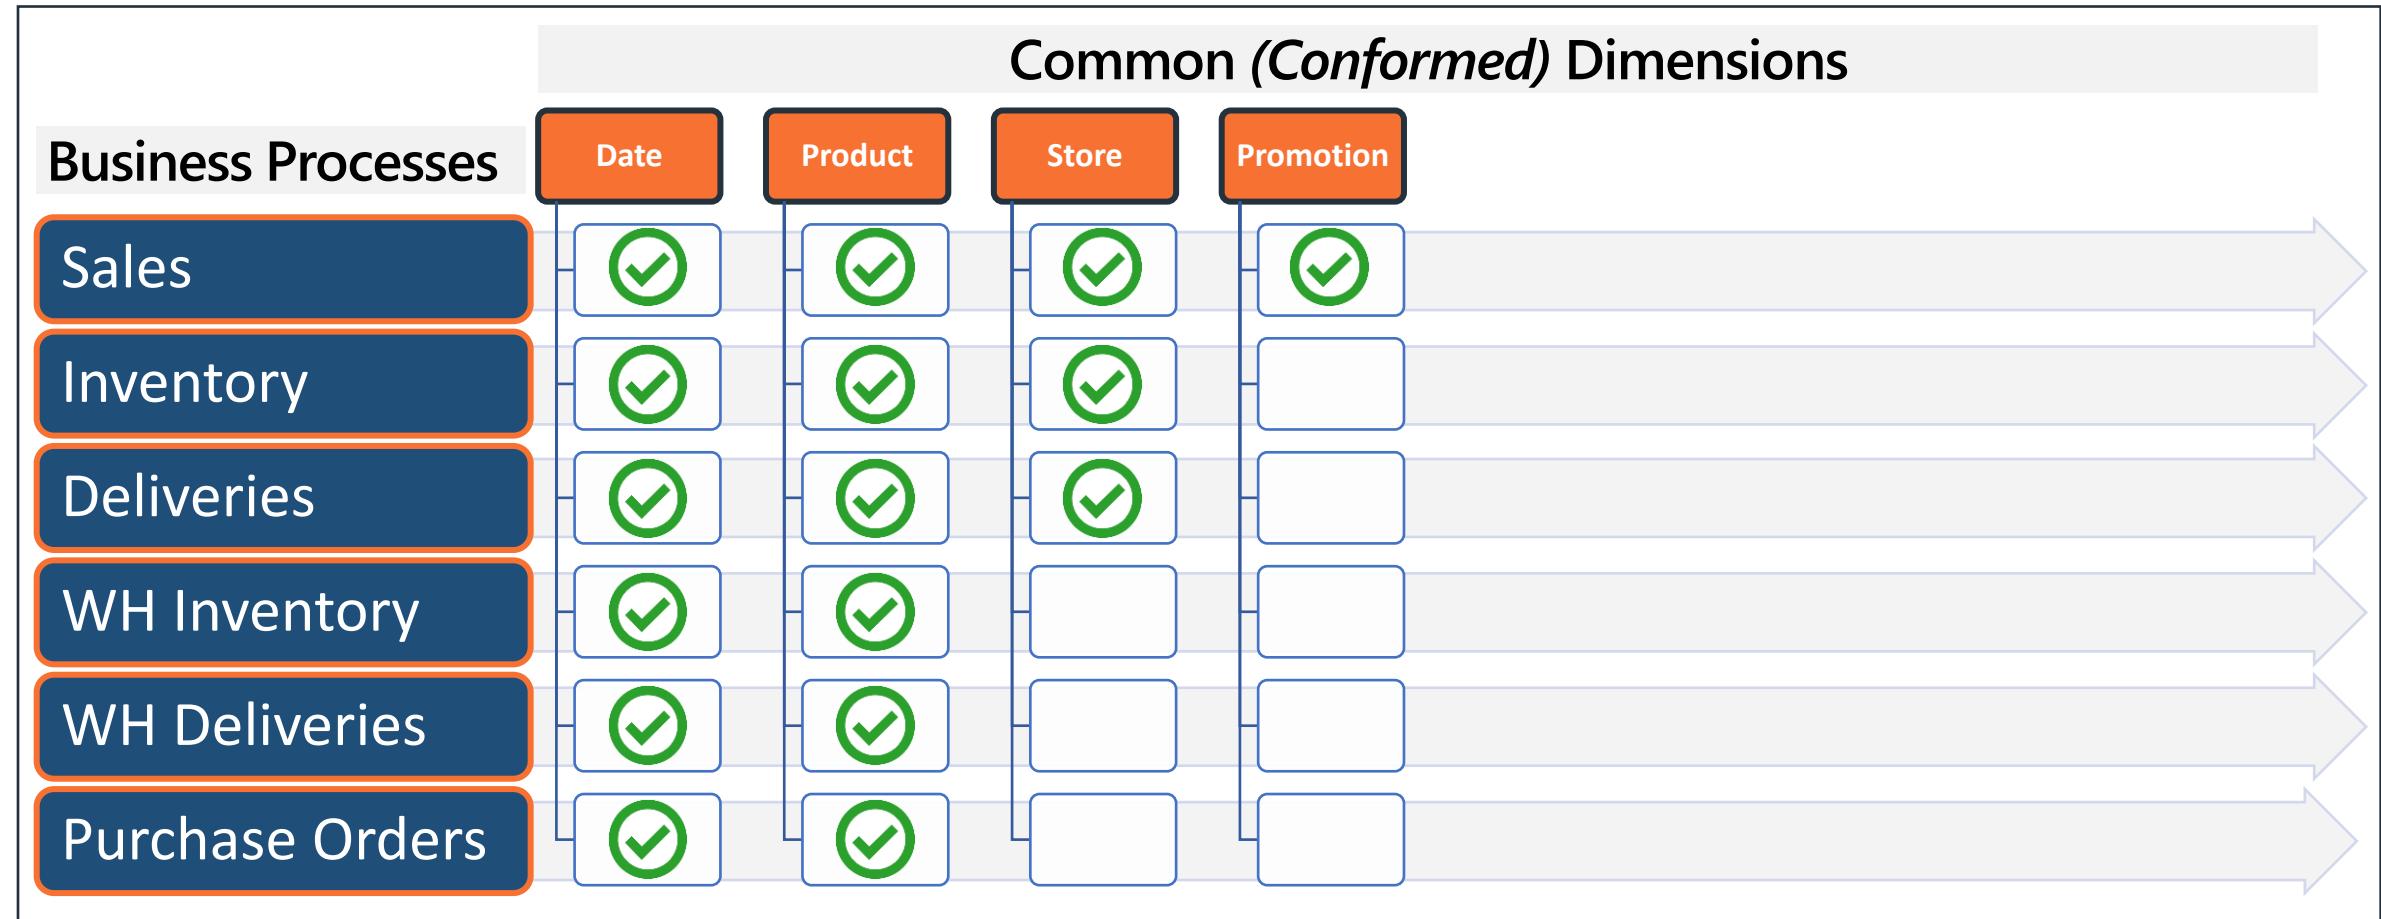


The Data Warehouse Toolkit, by Ralph Kimball



Identifying Dimensions: Data Warehouse Bus Matrix

Using the Same Dimensions Across Multiple Business Processes Enforces a Unified View of the Truth



The Data Warehouse Toolkit, by Ralph Kimball



Identifying Dimensions: Data Warehouse Bus Matrix

Using the Same Dimensions Across Multiple Business Processes Enforces a Unified View of the Truth

Business Processes	Common (<i>Conformed</i>) Dimensions							
	Date	Product	Store	Promotion	Warehouse	Vendor	Contract	Shipper
Sales	✓	✓	✓	✓				
Inventory	✓	✓	✓	✓				
Deliveries	✓	✓	✓	✓				
WH Inventory	✓	✓			✓	✓		
WH Deliveries	✓	✓			✓	✓		
Purchase Orders	✓	✓			✓	✓	✓	✓

The Data Warehouse Toolkit, by Ralph Kimball

Data Integration

How to Approach Populating a Data Warehouse



UNIVERSITY *of* VIRGINIA



Data Processing: Extract-Transform-Load (ETL)

Frequently, Data Must Be Moved from Sources to a Database and/or Data Lake



Extract

- This is the step where sensors wait for upstream data sources to land. Once available, we transport the data from their source locations to further transformations.



Transform

- The heart of any ETL job: apply business logic, perform actions such as filtering, grouping, and aggregation to translate raw data into analysis-ready datasets.



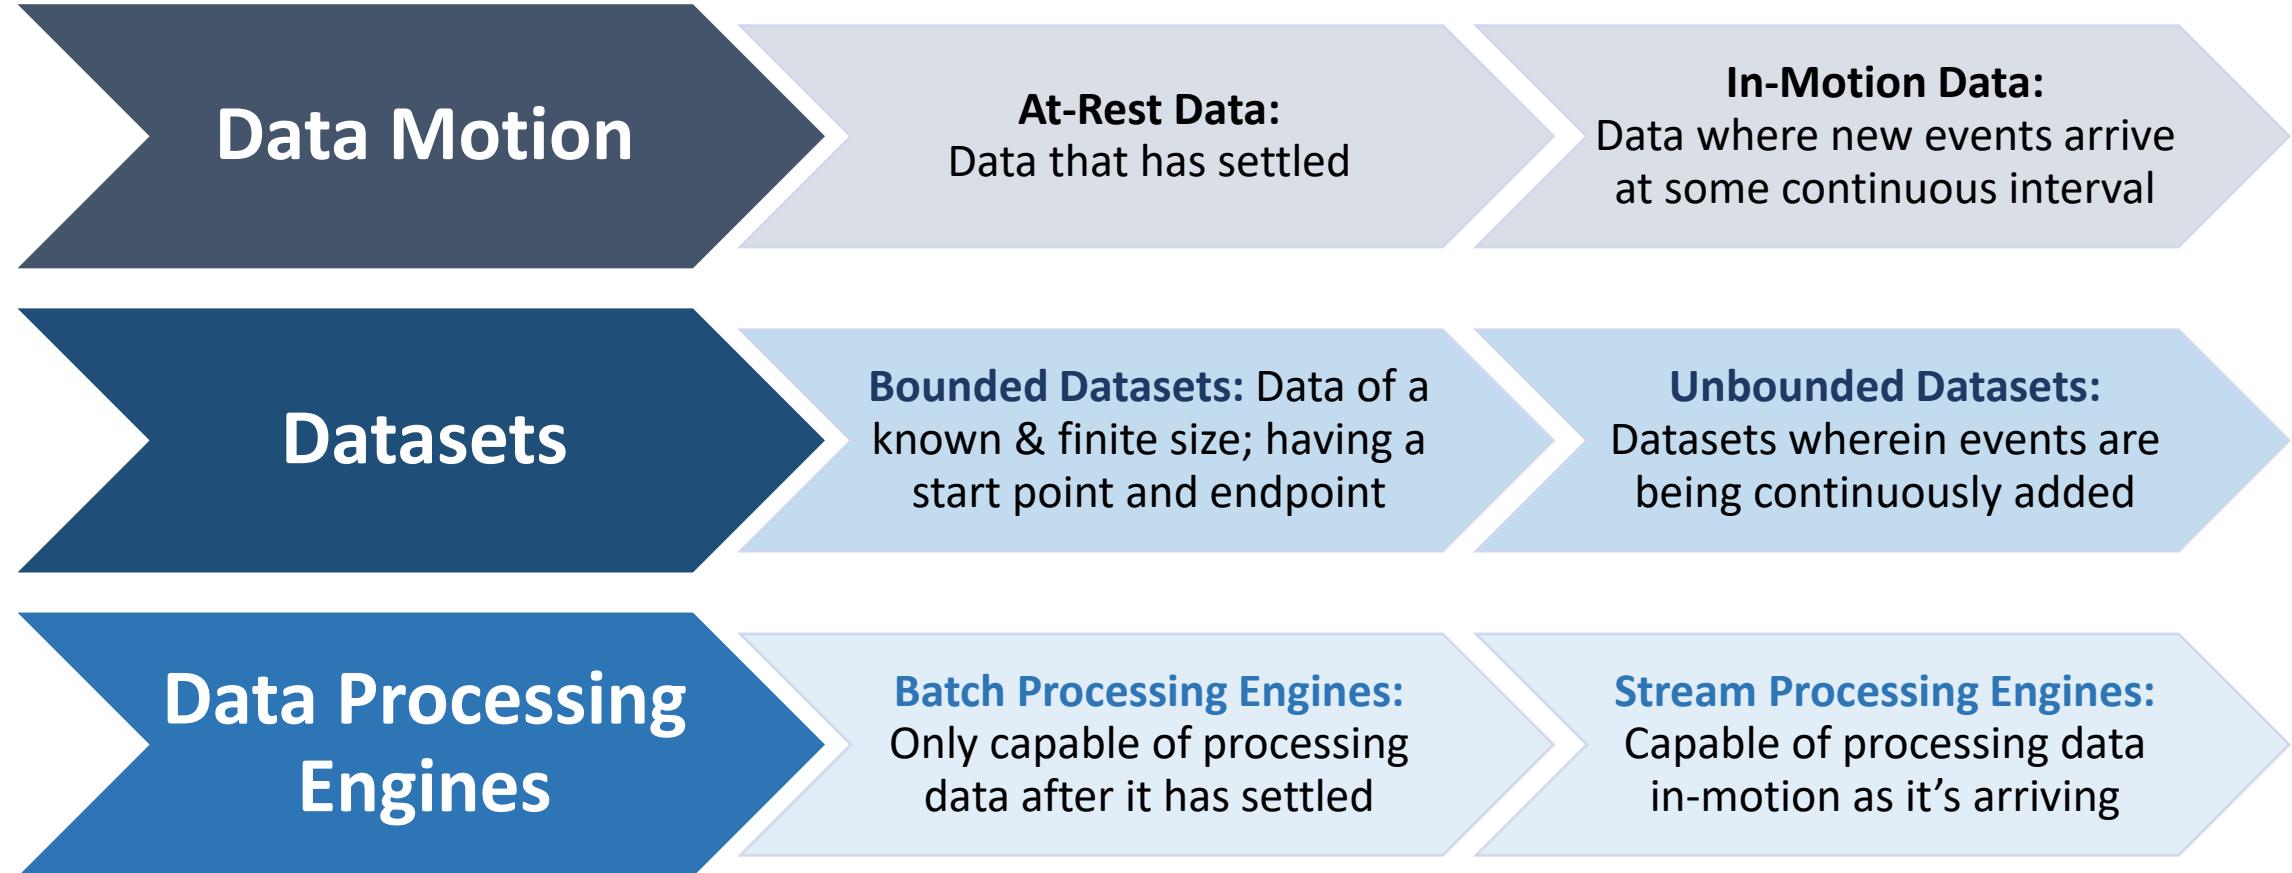
Load

- Load the processed data and transport to a final destination. Can now be consumed directly by end-users or treated as yet another upstream dependency.



Data Processing: Batching versus Streaming

The rate at which incoming data must be integrated dictates how it must be processed





Data Processing Paradigms: Latency Requirements

Latency & Response:

The speed at which clients require new insights determines the frequency at which new data must be processed

1. Batch

2. Continuous/Streaming

3. Real-time

10 ms

100 ms

1 sec

1 min

1 hour

1 day

Low-Latency Real-Time

- Spark-less, highly-available prediction server

Real-Time

- Prediction server with Spark

Micro-Batch

- Structured Streaming

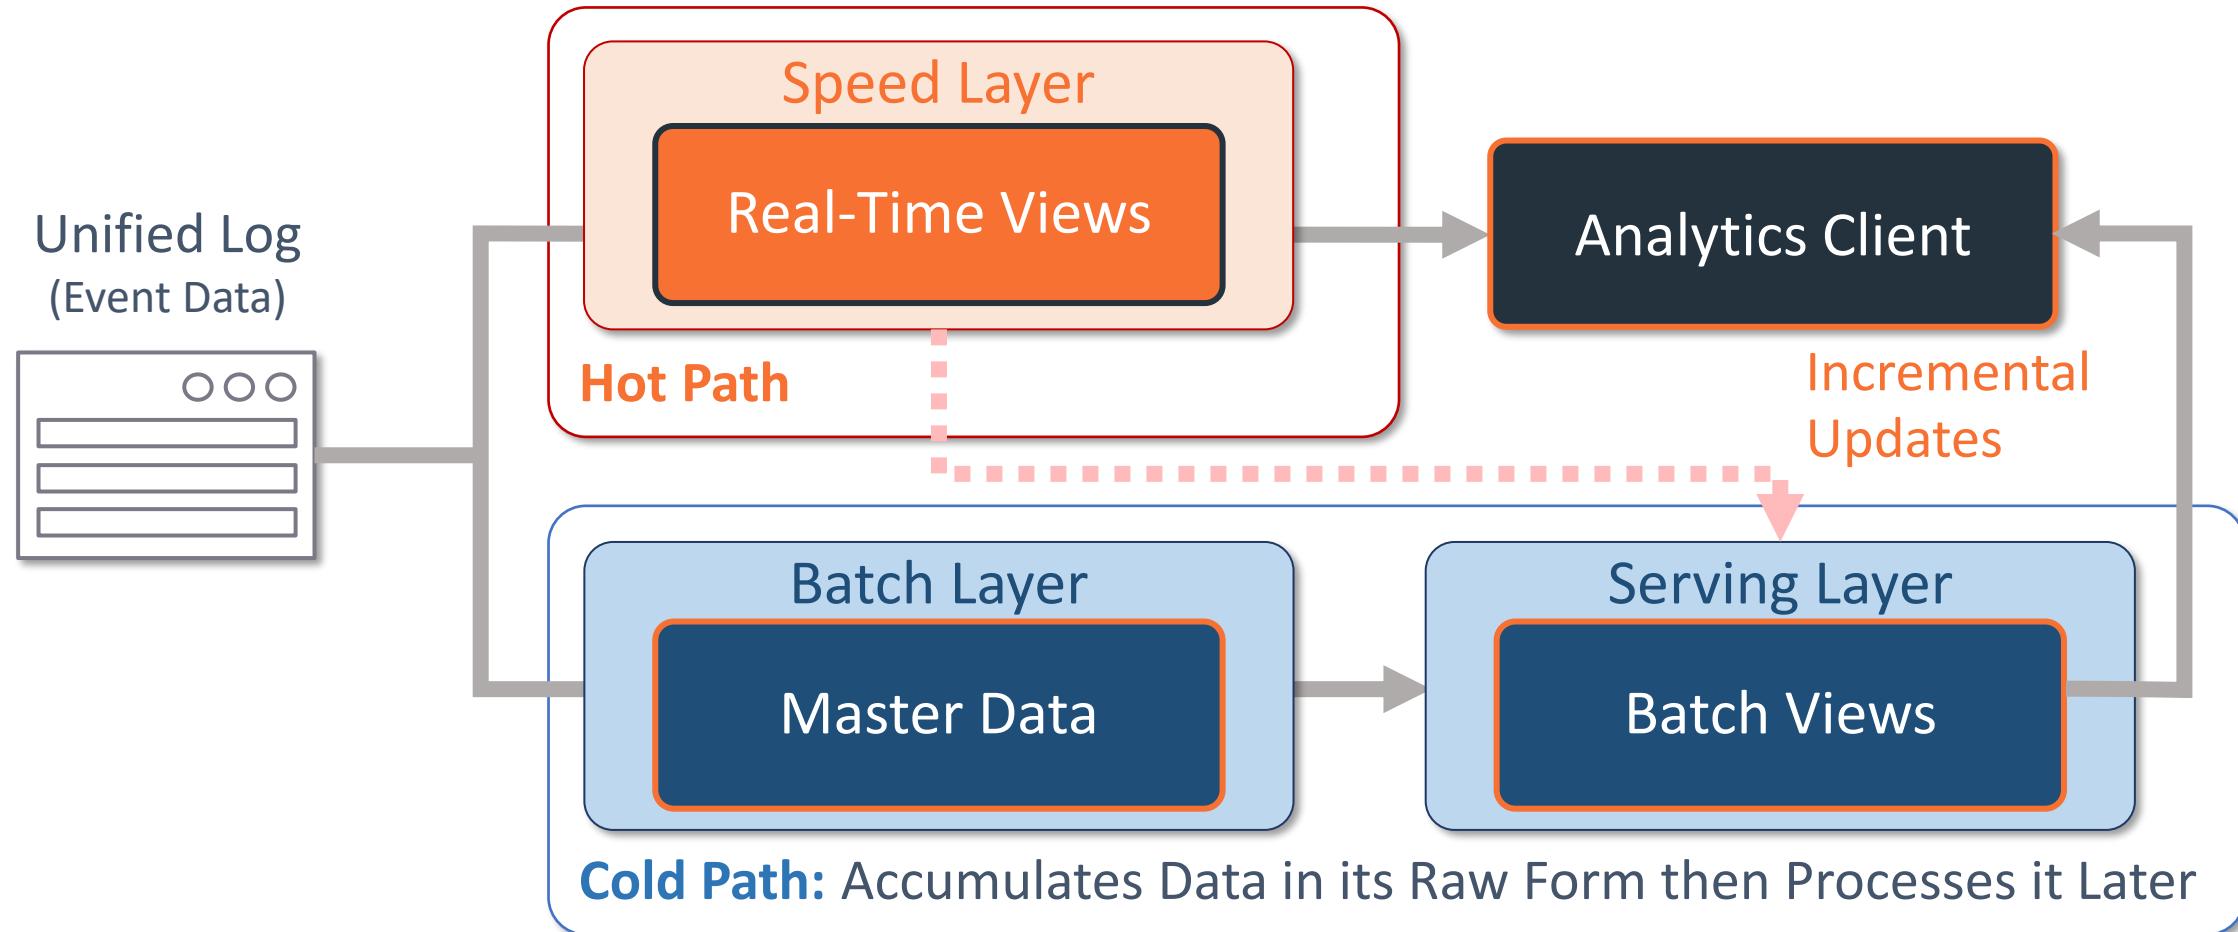
Batch

- Spark batch processing



Data Processing Paradigms: Lambda Architecture

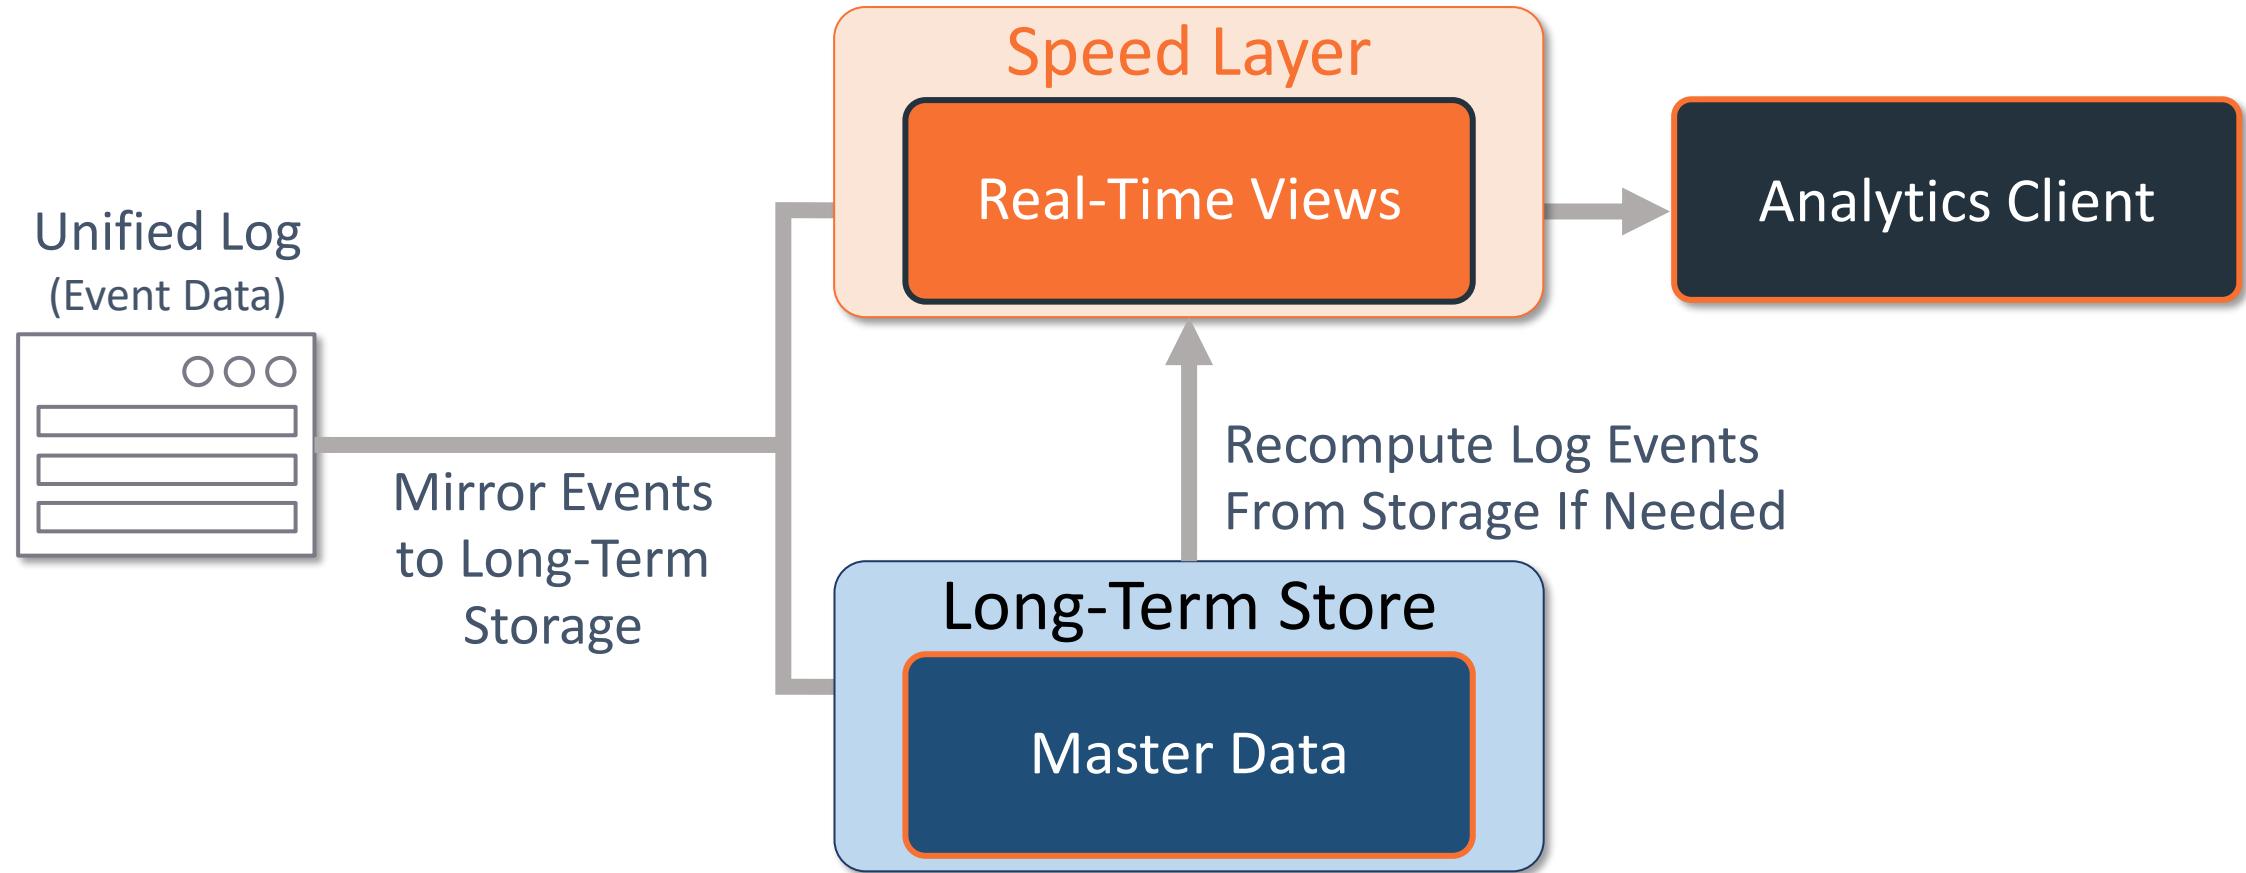
All Data Flows Through One of Two Paths: Hot or Cold





Data Processing Paradigms: Kappa Architecture

All Data Flows Through One of Two Paths: Hot or Cold





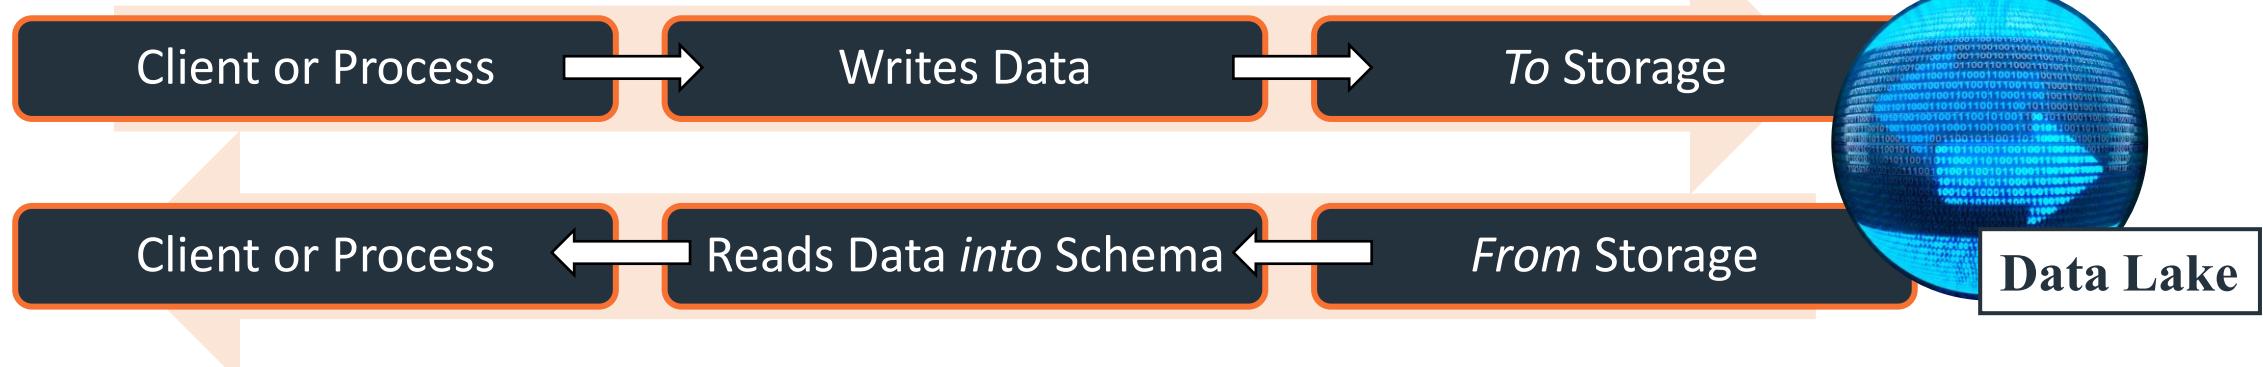
Paradigms: Data Storage and Retrieval

Schema on Write versus Schema on Read

Schema on Write: Stores data in an organized structure (schema)



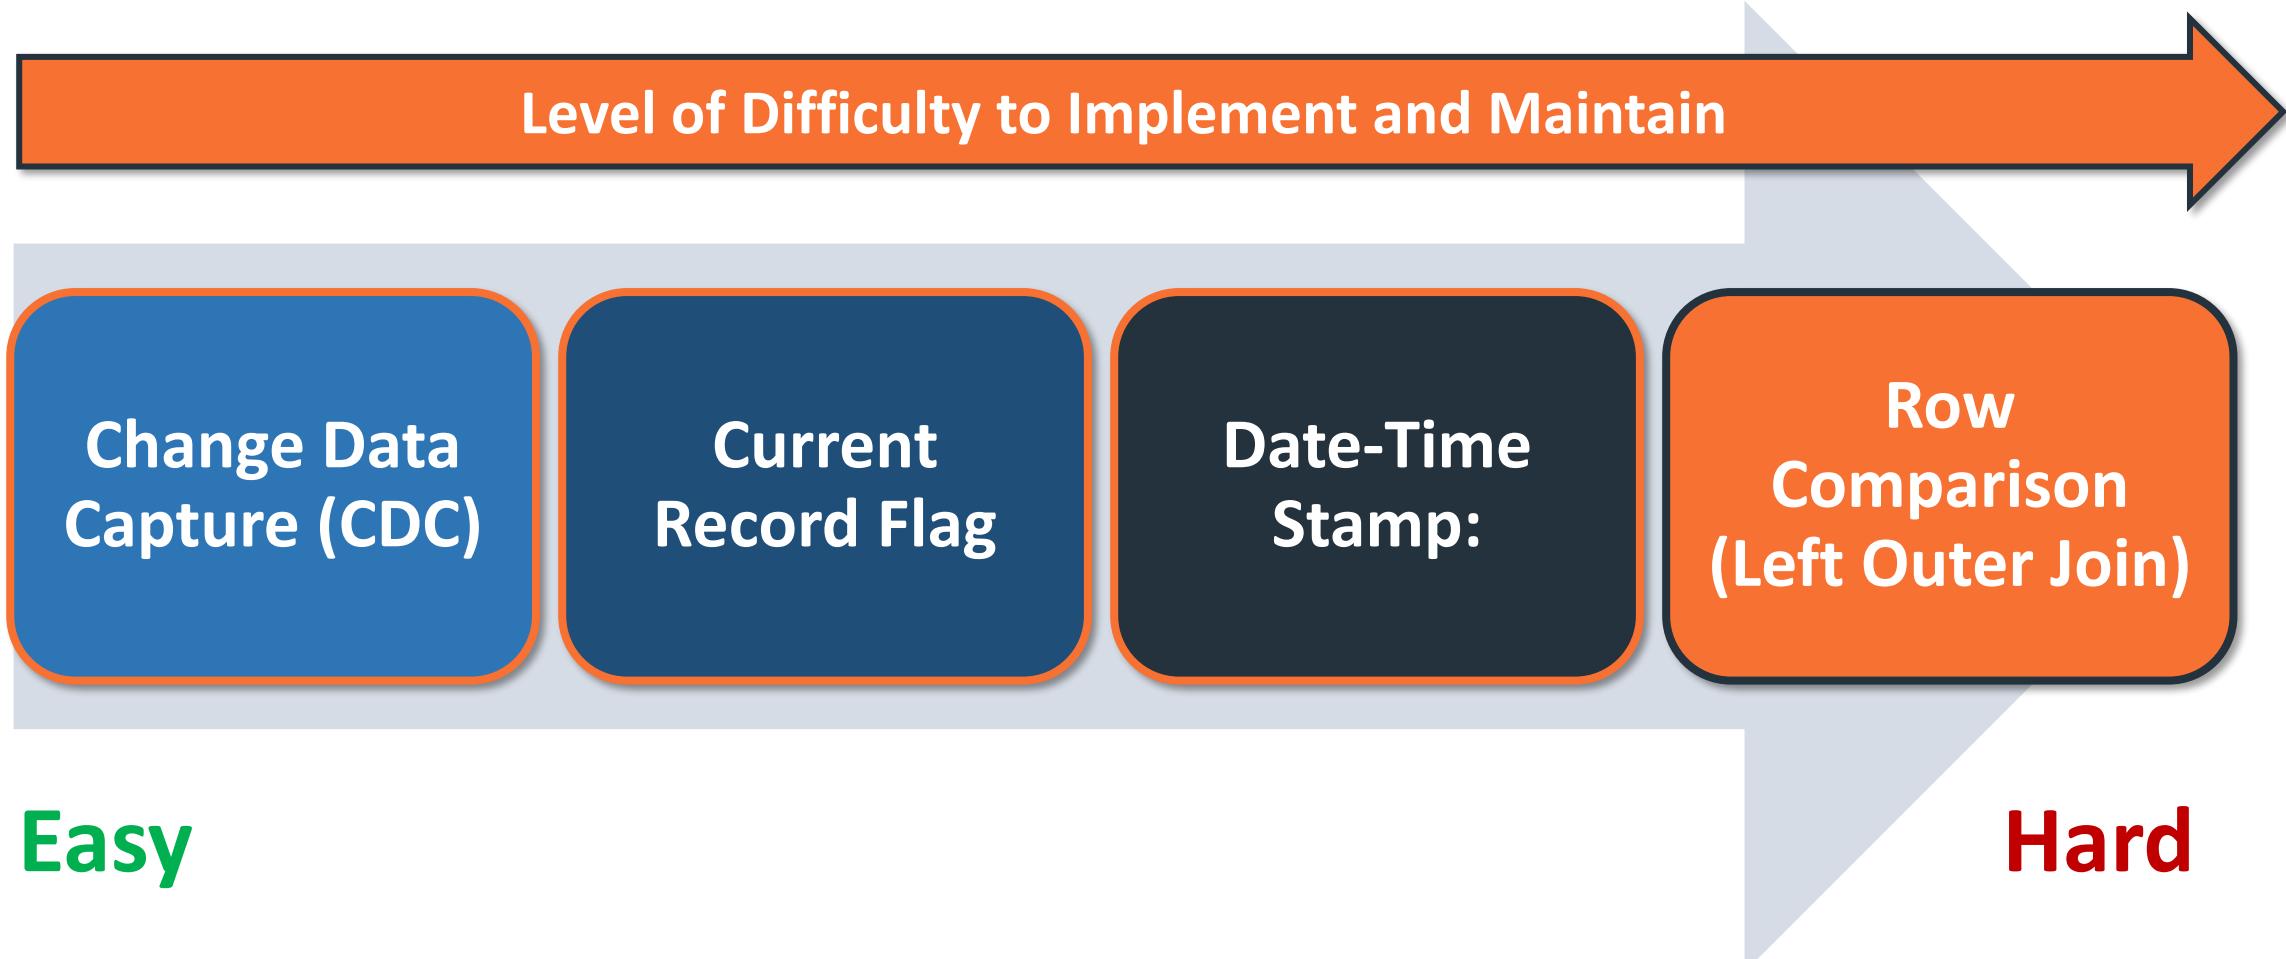
Schema on Read: Applies schema only when read, data stored in its original format





ETL Processing: Incremental Extraction

Techniques for Minimizing Data Movement: Extract Only the Changes





Data Integration Patterns: Dimensional Data

Slowly Changing Dimension Update Strategies: Handling Variable Rates of Change

SCD Type 0

- Data in the Column Never Changes: Ever!
- Only for Static Reference Data

SCD Type 1

- No History is Maintained
 - Existing Values are Overwritten by New Values
- **UPDATE**

SCD Type 2

- Historic Values are Maintained
 - New Values are Written to a New Row
 - *IsCurrent* Flag
- **INSERT**

SCD Type 3

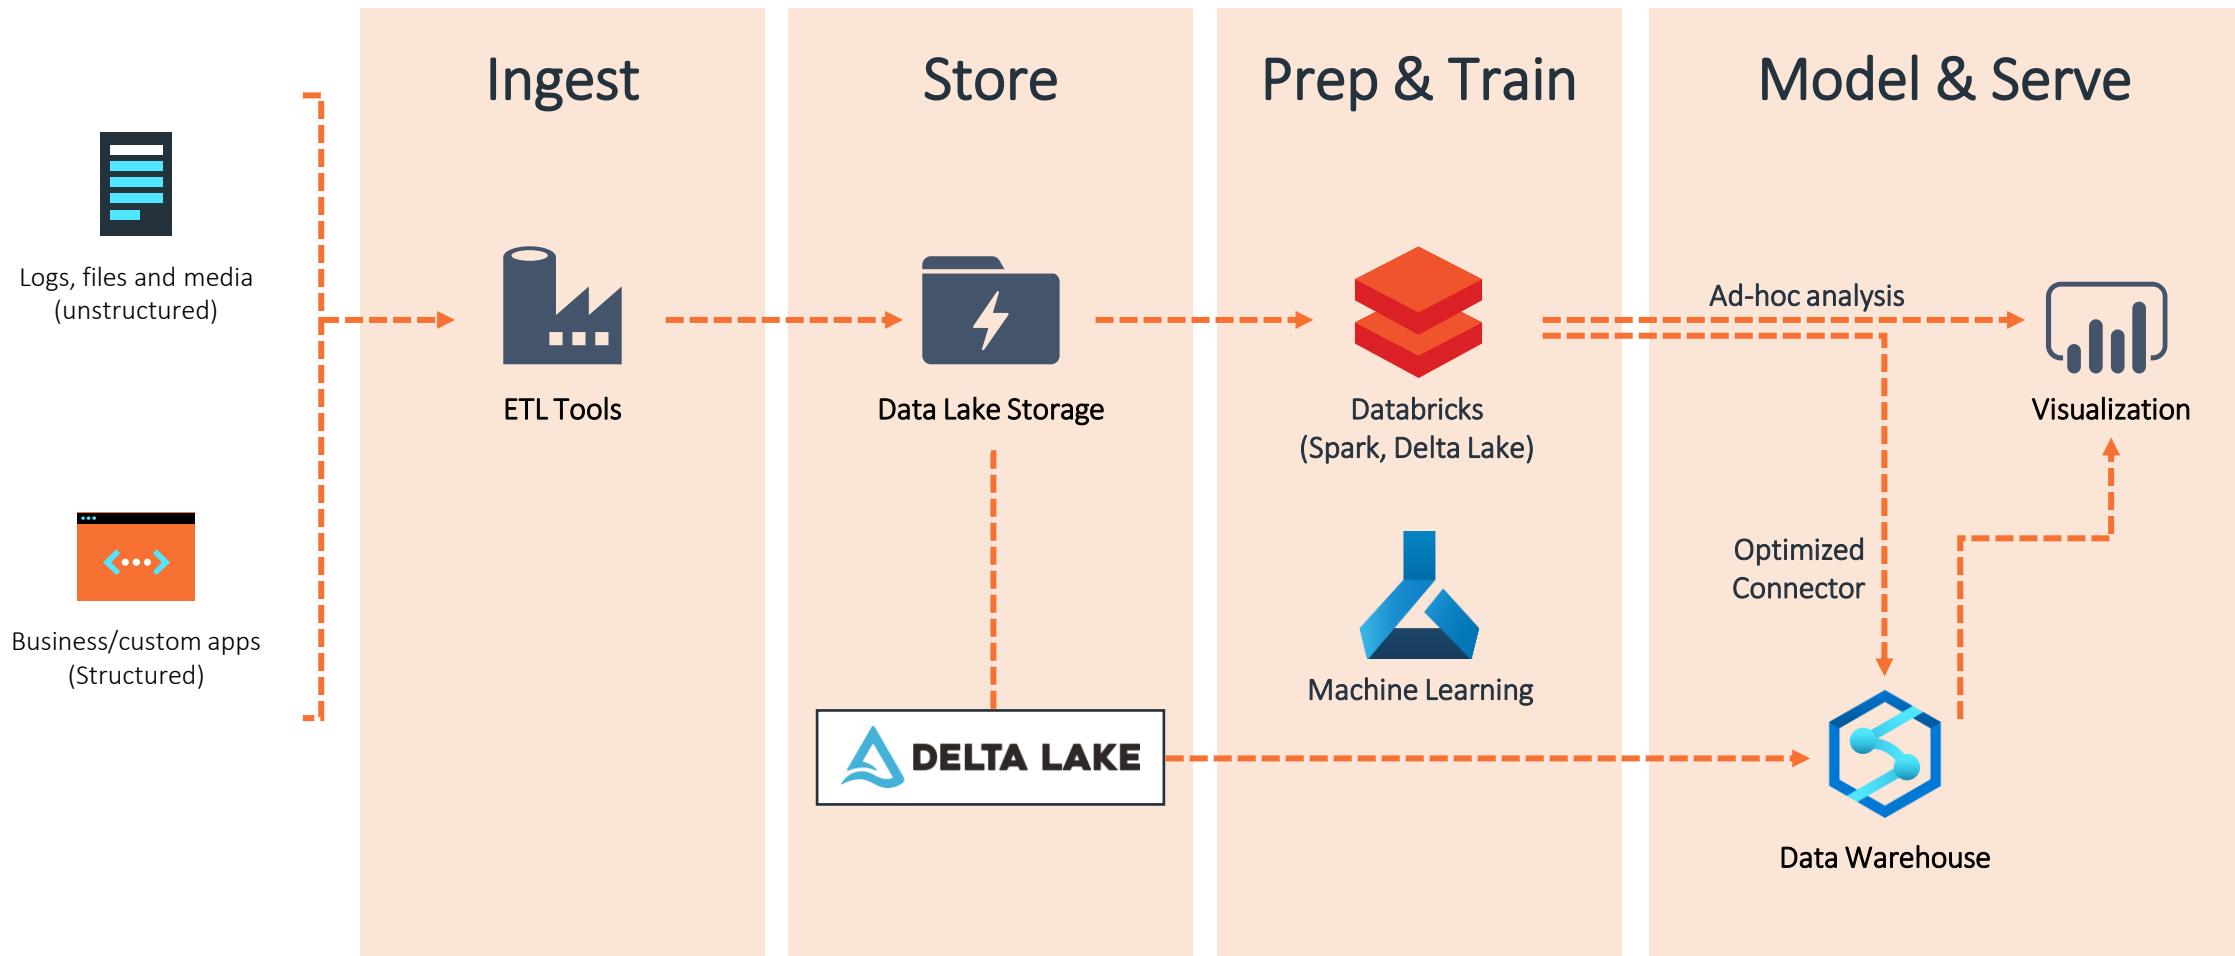
- A New Current Value Column is Created in the Existing Record
- Original Column is Also Retained

Easier to Implement and Maintain

More Difficult to Implement and Maintain

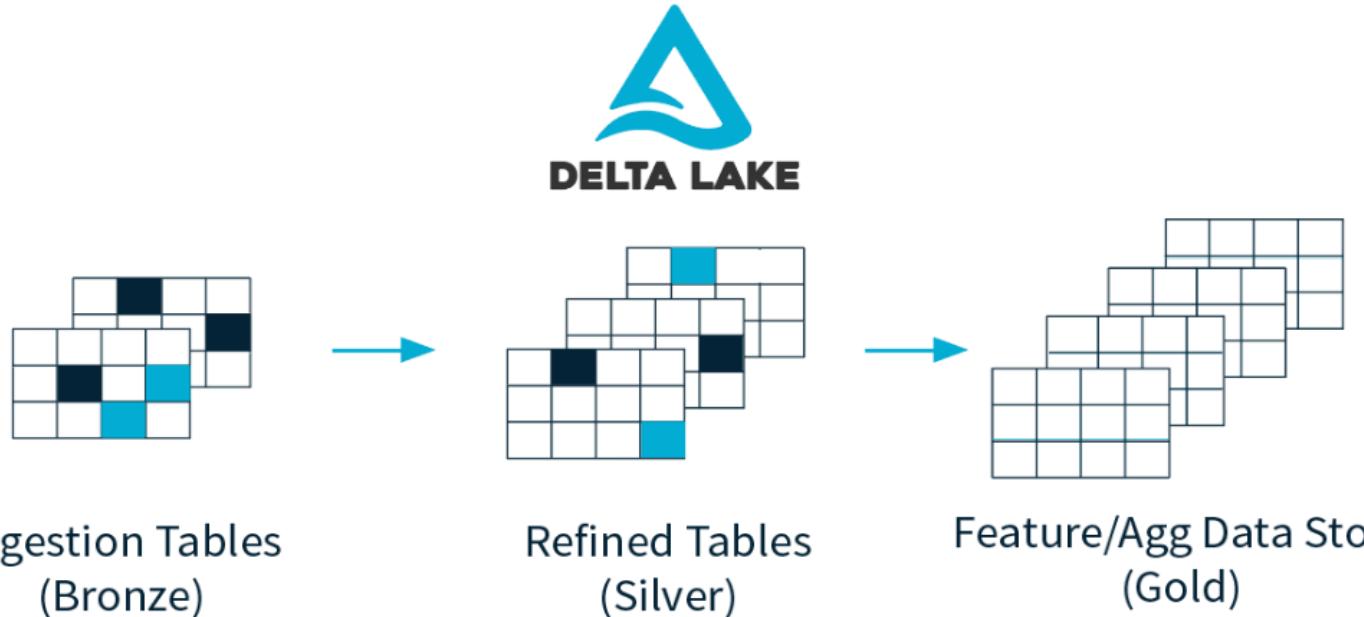


Data Engineering... for Data Science





Databricks: Delta Lake at Scale



ACID Transaction Guarantees

Atomic, Consistent,
Isolated, Durable

Versioned Parquet Files

Delta transaction log keeps
track of all operations

Efficient Upserts

MERGE, DELETE, UPDATE

Time Travel

Audit history, pipeline
debugging, data
reproducibility

Small file compaction with no interrupt to availability

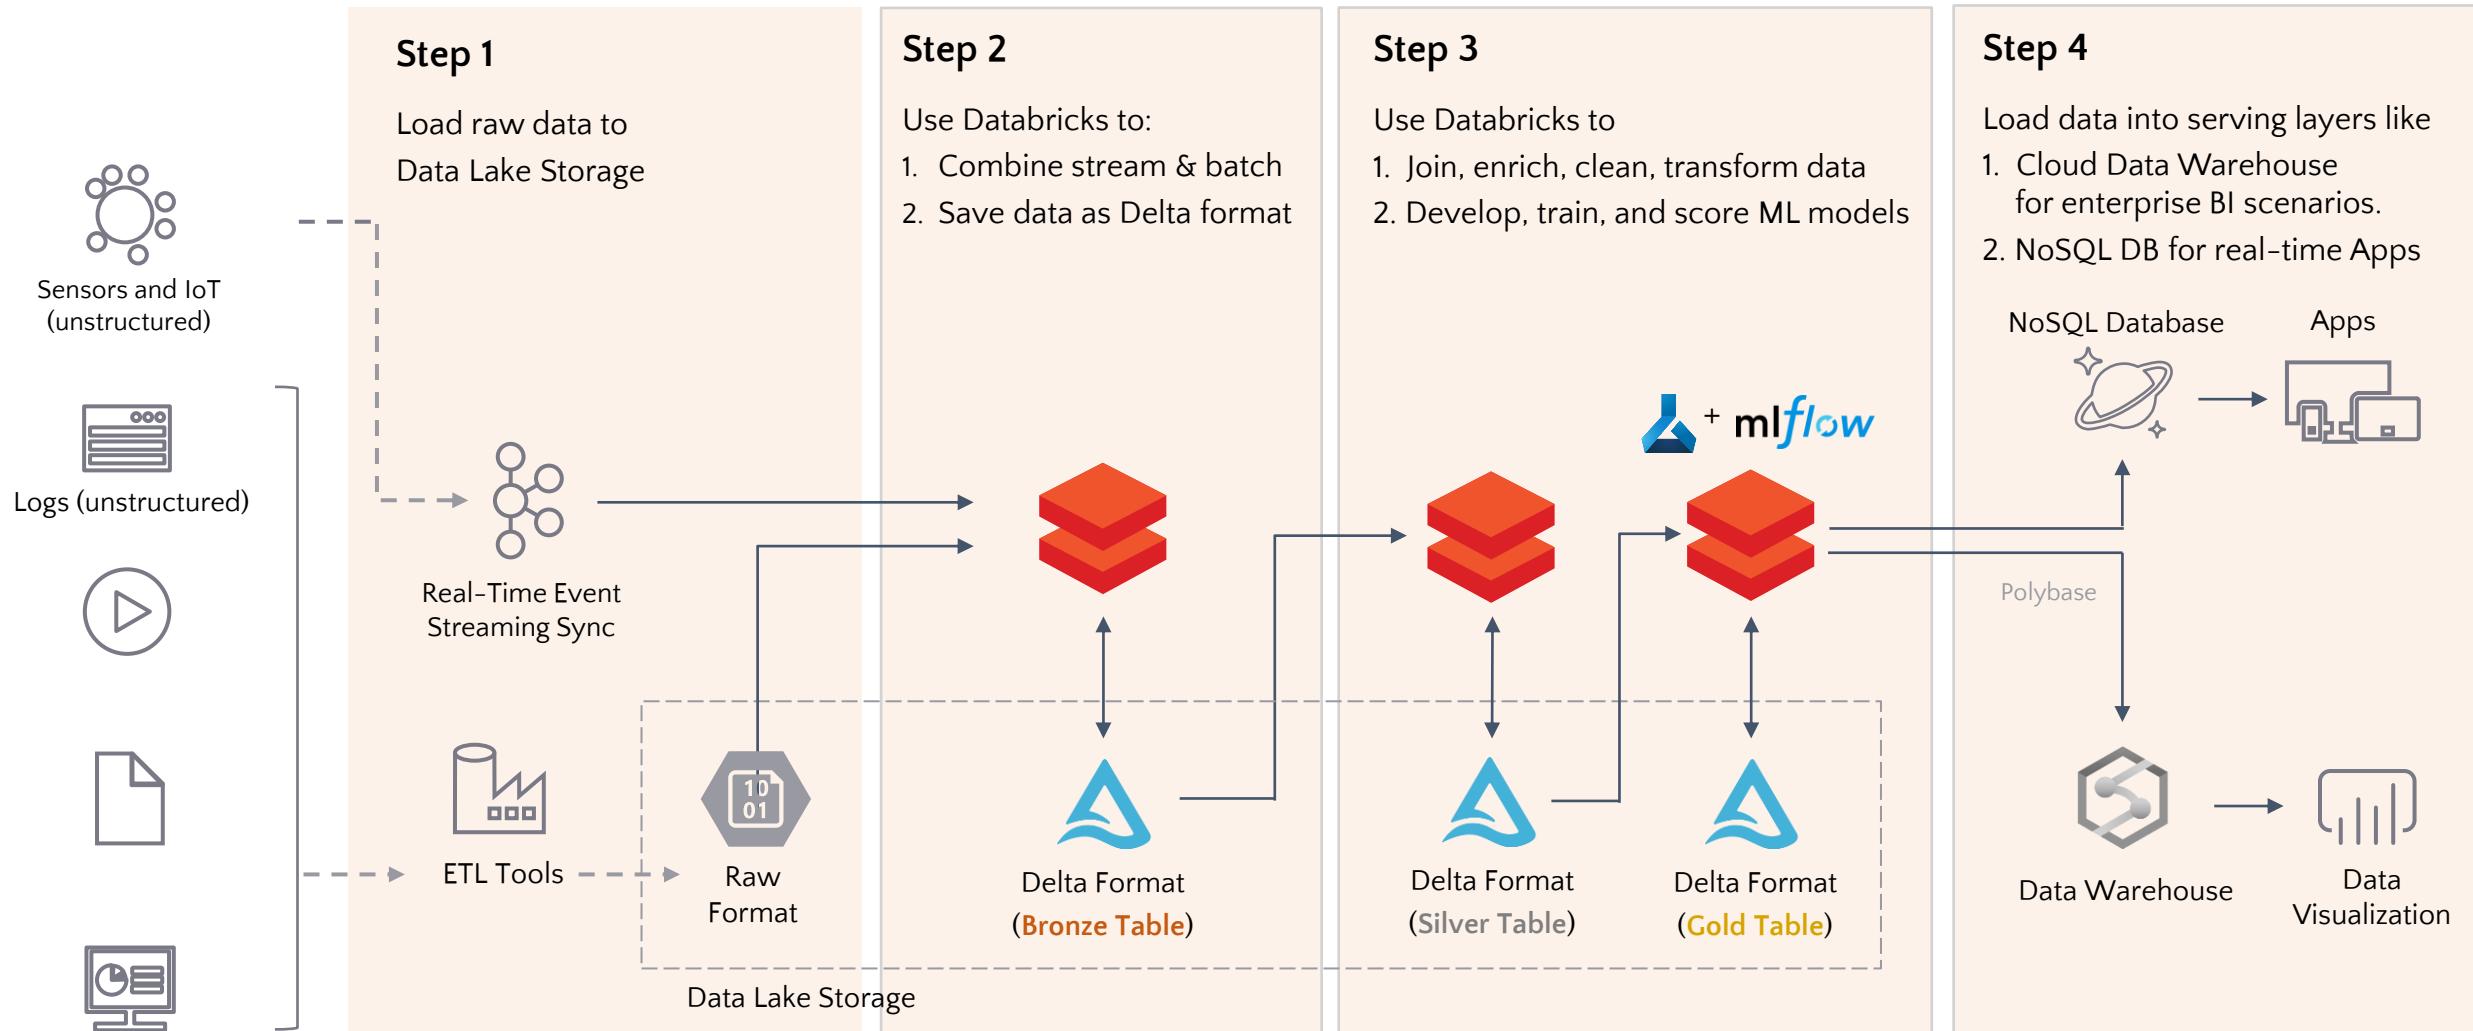
OPTIMIZE and VACUUM

Z-Order partitioning with up to 100x perf

New multidimensional partitioning enables
data skipping



Design Pattern: Modern Data Warehousing



Q & A

An Overview of Data Warehouse Systems